

tradeforme — Next Steps

Action list of things only **you** can do. Everything below is blocked on your input, your credentials, or your decision — not on more code.

Where the project stands

Repository	https://github.com/EtixP/tradeforme (private)
Branch	main
Milestones done	M1, M2, M3 (deterministic extraction, 96.6% OK), M4, M5 (signals stored, v2 with KOSP)
Tests passing	92 / 92
Disclosures in DB	514,304 across 488 trading days (Jun 25 2024 – Jun 24 2026, full 24 months)
Supply-contract events	3,531 candidates → 3,412 OK + 119 manual-review (96.6% / 3.4%), 0 blocked
Real contract values extracted	3,412 events with verified contract_value, prior_year_revenue, ratio
Candidate signals stored	693 (v2 strategy: ratio ≥ 0.08, KOSPI only, skip government counterparty). Mean strength
Price data source	pykrx (free, no auth)
LLM extraction	Code complete (prompts + validator + client) — optional now that deterministic parser wo
Broker integration	PaperBroker (historical fills) ready; KIS broker needs your credentials

Threshold sweep on 24 months / 3,412 events

Net of 0.313% roundtrip cost (commission + VAT + sale tax + slippage). Loop-3 ran the same sweep on only 514 events (3 months); the comparison column shows how those earlier numbers held up after 6.6× more data.

Threshold	n (24mo)	T+1 net	T+5 net	T+5 win%	PF (5d)	Median 5d	Loop-3 T+5 (n=514)
0.00 (all)	3,412	−0.12%	+0.35%	44.1%	1.21	−0.68%	+1.21% (n=514)
0.08 «	2,176	−0.07%	+0.22%	42.9%	1.16	−0.73%	+0.55% (n=333)
0.15	1,177	−0.16%	+0.56%	44.0%	1.25	−0.71%	+0.89% (n=169)
0.20	819	−0.06%	+0.59%	42.9%	1.25	−1.07%	+1.15% (n=110)
0.30	470	+0.18%	+0.50%	41.2%	1.21	−1.29%	+3.10% (n=59) «
0.50	217	+0.47%	+0.92%	38.1%	1.31	−1.79%	+2.67% (n=23)
1.00	66	−0.06%	−1.83%	32.3%	0.69	−3.98%	—

The big finding: the ≥ 0.30 "edge" from Loop 3 was largely noise. Going from 59 to 470 events, the T+5 net mean dropped from +3.10% to +0.50% and profit factor from 1.93 to 1.21. Median is negative at every threshold. The 1.00+ row is actually negative — suggests very large contracts may be a slight negative

signal (possibly representing distressed wins or unusual situations).

Honest conclusion: after 24 months and 3,412 events, the major-supply-contract event does *NOT* produce a clean tradable post-disclosure drift after costs. Mean net returns at every threshold are within noise of zero, win rates are 38–46%, medians are negative. This is exactly the kind of "strong negative result" CLAUDE.md anticipates — it answers the original research question. The market reacts efficiently to these disclosures by event-day close.

What can still rescue this: (a) sub-segment by counterparty type, sector, market cap, or recent volatility — maybe certain slices DO work. (b) try other event types (buybacks, dilutive financing, halt/resumption). (c) different exit horizons (T+10, T+20). (d) the M5b ML filter on richer features. Each is a real hypothesis to test — none is guaranteed to work.

Loop-6 subgroup + walk-forward results

Sliced the 3,377 events by market, ratio bucket, contract value, day-of-week, calendar quarter, and counterparty type (regex heuristic). Then walk-forward validated the most-positive subgroup and the most-negative subgroup across four 6-month windows. Honest reading below.

Subgroup	Aggregate T+5	Aggregate PF	Walk-fwd verdict (4 windows)
KOSPI / ratio 0.15-0.30	+1.59% (n=199)	1.59	ALTERNATES (+/-/-/+) — not robust
KOSPI (any ratio)	+0.63% (n=1708)	1.22	3/4 windows positive — modest but holds
KOSDAQ (any ratio)	+0.06% (n=1669)	1.02	no edge
Government counterparty	−2.14% (n=90)	0.51	3/4 windows NEGATIVE (sometimes −5%)
Other_korean_corp	−0.65% (n=417)	0.84	consistently negative
Ratio 1.00+	−1.83% (n=66)	0.64	small n; likely distressed/unusual events

Two actionable findings:

- **Government-counterparty contracts consistently underperform.** 3/4 walk-forward windows show net T+5 returns of –5% to –0.5%. Aligns with intuition: government contracts are pre-priced, low-margin, expected. We can't short (CLAUDE.md rule), but this is a **strong negative filter** — the strategy should skip long signals where counterparty is government.
- **KOSPI > KOSDAQ** in 3/4 windows. Aligns with the liquidity-premium hypothesis. Worth restricting the strategy to KOSPI-only.

Two findings ruled out: the KOSPI/0.15-0.30 subgroup that looked best in aggregate (+1.59% T+5) alternates positive/negative across windows — a classic small-sample fluke. The 0.30+ ratio "edge" from Loop 3 stays dead at scale. Median return remains negative at every threshold.

Raw CSV: data/subgroup_analysis.csv (3,377 events x 6+ feature columns).

Raw CSV: data/event_study_results.csv (3,531 events, all horizons, gross). Loop-3 log of the smaller dataset is preserved in git history.

Action items

Sorted by priority. P1 unblocks the next milestone. P2 hardens the existing system. P3 is optional polish.

Priority	What	Why	How
P1	Decide on the strategic pivot	The 24-month backfill killed the app	Reply 30% the old path, loop 70% the new one. Supply can be tested
P1	Get an LLM API key (Anthropic vs OpenAI)	LLM needed for ML initiative	Go to OpenAI (too expensive) or try a different language to strengthen the AI
P2	Rotate the DART API key	The current key briefly appeared in chatGPT during the 1111 Smokey test (1111 is the replacement)	
P2	Set up GitHub branch protection	Once live trading code exists, you don't want to lose it	GitHub Actions: Stay Put before main -> Settings -> Branches -> add
P2	Back up data/kdtb.db	111k rows of ingested DART data	Easy to replicate into a USB drive or backup to the Desktop or data
P2	Schedule daily DART ingestion	Right now ingestion only runs when you are home	Use the CLI manually/Desktop cron job to schedule
P2	Obtain KIS broker credentials	Required for Milestone 6 (broker integration)	Go to the investor's securities/live
P2	Decide whether to raise min_confidence threshold	Default is 0.9, but the sweep requires 0.95	0.95 is a good middle ground
P2	Review risk thresholds in config/default.yaml	Default is 0.3, but the sweep requires 0.35	0.35 is a good middle ground
P3	Install gitleaks pre-commit hook	Catches an accidental commit of .env	brew install gitleaks & gitleaks pre-commit -B
P3	Decide if PDF action lists should be included	Currently, NEXT_STEPS.pdf will be included	local only as per NEXT_STEPS work to mitigate a potential
P3	Decide on Korean market holiday handling	Right now, the ingestion CLI runs every day	Optional add a holiday calendar (only for holidays, business)

Where to find things

Project root	/Users/jsp2022310/Desktop/tradeforme
Repo on GitHub	https://github.com/EtixP/tradeforme
Project spec	CLAUDE.md (in repo root)
Source code	src/kdtb/ — schemas, config, data, strategy, risk, storage
Tests	tests/ — run with: <code>.venv/bin/pytest</code>
DART ingestion CLI	<code>scripts/ingest_disclosures.py --date YYYY-MM-DD</code>
Event study CLI	<code>scripts/run_event_study.py [--limit N]</code>
Parse docs CLI	<code>scripts/parse_supply_contracts.py [--limit N] [--skip-existing]</code>
Filtered event study	<code>scripts/run_filtered_event_study.py [--min-ratio 0.08]</code>
Strategy runner	<code>scripts/run_strategy.py [--min-ratio 0.30] [--dry-run]</code>
DB backup	<code>./scripts/backup_db.sh</code> — gzipped SQL dump under <code>data/backups/</code>
Walk-forward (NEW)	<code>scripts/walk_forward.py</code> — reproduces the v0/v1/v2 4-window table
Local DB (gitignored)	<code>data/kdtb.db</code> — SQLite, inspect with <code>sqlite3</code> or DB Browser
Event study CSV	<code>data/event_study_results.csv</code>
Secrets (gitignored)	<code>.env</code> — <code>DART_API_KEY</code> lives here
Config	<code>config/default.yaml</code> , <code>config/paper.yaml</code> , <code>config/live_tiny.yaml</code>
Live-trading guard	<code>src/kdtb/config.py</code> — raises if YAML says LIVE without <code>ENABLE_LIVE_TRADING=true</code> in

How to verify the project is healthy

Run these from the repo root any time you want to confirm nothing is broken:

```
.venv/bin/pytest
```

```
.venv/bin/python -m kdtb.main
```

```
.venv/bin/python scripts/ingest_disclosures.py --date $(date +%Y-%m-%d)
```

What Claude built across the autonomous loops

Loop 1 (M4 + M5):

- **data/market_data_client.py** — pykrx wrapper for OHLCV + event-time returns
- **scripts/run_event_study.py** — batched event study CLI
- **strategy/major_supply_contract.py** — M5 strategy engine (was placeholder)

Loop 2 (cost model + M3 prep + paper broker):

- **backtest/cost_model.py** — Korean equity costs
- **backtest/metrics.py** — per-trade metrics
- **interpretation/llm_prompts.py, extraction_validator.py, llm_client.py**
- **broker/base.py + broker/paper_broker.py**

Loop 8 (walk-forward validation of v0/v1/v2):

The v2 aggregate looked promising; this loop verified that the improvement holds across 4 non-overlapping 6-month windows (the real test of "edge vs curve-fit").

Window	v0 T+5 / PF (n)	v1 T+5 / PF (n)	v2 T+5 / PF (n)
Jul24-Dec24	+0.04% / 1.01 (785)	−0.38% / 0.90 (498)	+1.57% / 1.58 (179) «
Jan25-Jun25	+0.59% / 1.23 (713)	+0.51% / 1.19 (463)	+1.11% / 1.65 (138) «
Jul25-Dec25	+0.41% / 1.15 (930)	+0.67% / 1.25 (573)	−0.55% / 0.79 (200)
Jan26-Jun26	+0.45% / 1.10 (906)	+0.15% / 1.03 (591)	+1.01% / 1.28 (163) «
Aggregate	+0.37% / 1.11 (3,334)	+0.25% / 1.07 (2,125)	+0.72% / 1.27 (680)
Positive windows	4/4	3/4	3/4

Verdict: v2 holds up. 3 of 4 windows positive, average +1.23% in the three winning windows, −0.55% in the losing one. The losers are smaller in magnitude than the winners. Aggregate T+5 net mean +0.72%, profit factor 1.27, win rate 47.1%. v1 is consistently the worst of the three — the ratio filter alone removes signal without adding edge. The KOSPI + skip-government conjunction is what's doing the work.

Still not a strategy ready for real money. 47% win rate, one losing window in four, PF 1.27 (■1.27 won per ■1 lost) — thin margin once execution slippage drifts the numbers. But it IS the first thing in this project that survives walk-forward validation. The next moves (paper broker simulation, then ML on top) are now standing on real ground.

Reproduce with: `.venv/bin/python scripts/walk_forward.py`

Loop 5 (24-month backfill + scale-out + the negative finding):

- CLAUDE.md updated: ML/LLM allowed for extraction + feature enrichment + filtering, but NOT in the order-placing decision. New **Milestone 5b** added (ML-enriched signal filter).
- **scripts/backup_db.sh** — gzipped SQL dumps to data/backups/, keeps last 7, ready for cron.
- **24-month DART backfill:** 111,622 → **514,304 disclosures** across 488 trading days (Jun 2024 – Jun 2026).
- **Re-ran deterministic parser:** 533 → **3,412 OK extractions** (96.6% success rate held up at scale, 0 blocked).
- **Re-ran event study + threshold sweep:** the 30%+ edge from Loop 3 (T+5 +3.10%, PF 1.93) collapsed to T+5 +0.50% / PF 1.21 on 8× more data. Median negative at every threshold.
- **Event study script now checkpoints CSV every 100 stocks** — one of yesterday's runs hung overnight; we no longer lose hours of fetches.

Suite: 28 → 43 → 68 → 84 → **88** passing tests. Zero regressions across five loops.

Honest caveats

- The naive event study uses title-only filtering. It does not separate small contracts from large ones, so noise dominates. M3 (LLM extraction) is needed to filter by contract/revenue ratio.
- The event date used is the DART receipt date (YYYYMMDD). Intraday timing of the disclosure is not captured. If a disclosure is filed at 14:00 KST, our 'T+1' return actually includes the price reaction from 14:00 to next-day close, not a clean next-day-only return.

- Returns are raw close-to-close, not benchmark-adjusted. A market-wide rally on event days would inflate the mean.
- pykrx scrapes the KRX website. It's free but can be rate-limited. We sleep 0.15s between calls.
- LIVE_TINY broker code is not yet implemented. Even with KIS credentials, no real orders can be placed yet.